

Where is our AI token spend actually going?

Tokenomics Cost Intelligence: a dashboard for tracking LLM spend across publishers, models, applications and teams, and for measuring how much of that spend can actually be attributed.

4

LLM publishers

5

applications

13

weeks

2

role-based views

Hari Muppa

Swipe →

Four numbers an executive sponsor checks first


13-week totals. Week-over-week compares the latest week (W13) with the prior week (W12).

TOTAL TOKEN SPEND \$815.5K All publishers, all divisions ↑ 10.4% WoW	APPLICATION COVERAGE 84.5% Spend with app-level attribution ↓ 0.1 pts WoW
AMBIGUOUS COST \$126.2K Spend without full attribution ↑ 11.2% WoW	COST PER 1K TOKENS \$0.0130 Blended rate across all models ↑ 8.8% WoW

Roughly **\$1 in every \$6.50** of token spend could not be tied to a specific application.

Attribution erodes at every level you drill into


Share of total spend that can be attributed at each level of detail.



Almost all spend maps to a division (96.6%), but only **68.9%** maps to a specific user. That gap is where chargeback conversations break down.

How trustworthy is the number?


Every dollar is tagged by how its attribution was derived.



Only **60%** of spend is directly measured. The other 40% is estimated, modeled or unknown, so the dashboard shows data quality next to the dollars instead of hiding it.

Ambiguous cost spiked, then recovered


Weekly ambiguous cost % with a rolling 3-week average.



The 3-week average peaked at **21.0%** in W4 and has settled near **12.0%**. The rolling line keeps one noisy week from triggering a false alarm.

Spend is spread fairly evenly across divisions


13-week token spend by division.



The top five divisions sit within about **\$23.6K** of each other. With no obvious outlier, cost control has to come from how tokens are used, not from which team uses them.

Not all tokens cost the same

Cost per 1K tokens by token type, with each type's share of total spend.




Fine-tune tokens cost **2.7x** more than cached tokens and make up the largest share of spend. Publisher rates in this data differ by under 5%, so caching and fine-tune governance are the bigger levers.

Attribution quality by application

The detail page, with slicers for division, week, publisher, model, token type, tier, environment and audience.

Measured Estimated Modeled Unknown



Named applications are **61 to 79%** measured. The unattributed bucket is only **8%** measured and **48%** modeled, so the spend we can least explain is also the spend we trust least.

How this helps a business track token spend

1

Showback and chargeback

Spend by division, application and user gives finance a defensible basis for allocating AI costs to the teams that drive them.

2

Close attribution gaps

The coverage cascade shows exactly where tagging breaks down, so engineering knows which apps and users to instrument next.

3

Control unit cost

Cost per 1K tokens by token type, model and tier points to the real savings levers: caching, fine-tune governance and model routing.

4

Trust the numbers

Every dollar is labeled measured, estimated, modeled or unknown, so leaders know how much confidence to place in each figure.

Two views, one model: an Executive Summary for sponsors and a FinOps Analyst page with 8 slicers for drill-down.

TAKEAWAYS

Attribution is the real FinOps problem.

Track coverage, not just cost: division → app → user.

Show data quality next to the dollars.

Token type moves unit cost far more than publisher choice.

Give executives the headline and analysts the drill-down.

How is your team attributing LLM spend today?

Hari Muppa

harimuppagith.github.io